

Convolutional Neural Networks IV

CS7150, Spring 2025

Prof. Huaizu Jiang

Northeastern University

Recap

Convolution Summary

Input: $C_{in} \times H \times W$

Hyperparameters:

- **Kernel size:** $K_H \times K_W$
- **Number filters:** C_{out}
- **Padding:** P
- **Stride:** S

Weight matrix: $C_{out} \times C_{in} \times K_H \times K_W$
giving C_{out} filters of size $C_{in} \times K_H \times K_W$

Bias vector: C_{out}

Output size: $C_{out} \times H' \times W'$ where:

- $H' = \text{Ceil}((H - K + 2P + 1) / S)$
- $W' = \text{Ceil}((W - K + 2P + 1) / S)$

Common settings:

$K_H = K_W$ (Small square filters)

$P = (K - 1) / 2$ ("Same" padding)

$C_{in}, C_{out} = 32, 64, 128, 256$ (powers of 2)

$K = 3, P = 1, S = 1$ (3x3 conv)

$K = 5, P = 2, S = 1$ (5x5 conv)

$K = 1, P = 0, S = 1$ (1x1 conv)


$K = 3, P = 1, S = 2$ (Downsample by 2)

Properties: Shift-Invariant


Assume: I image, f filter

Shift-invariant: $\text{shift}(\text{apply}(I, f)) = \text{apply}(\text{shift}(I, f))$

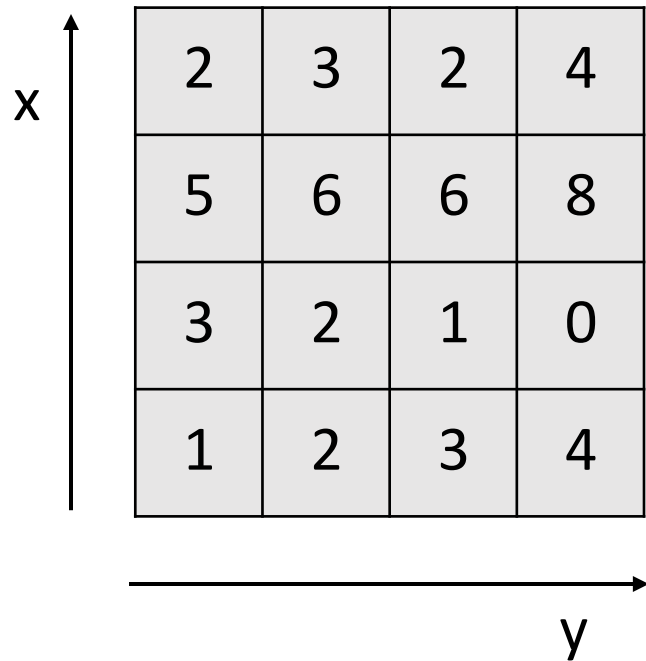
Intuitively: only depends on filter neighborhood



Max Pooling



Global Average Pooling



A 4x4 grid of numbers representing a single channel of a feature map. The grid is labeled with 'x' for the vertical axis and 'y' for the horizontal axis. The values in the grid are:

2	3	2	4
5	6	6	8
3	2	1	0
1	2	3	4




3.25

Average per channel (1 channel here).

Gradients?

Example: Lenet-5

Layer	Output Size	Weight Size
Input	1 x 28 x 28	
Conv ($C_{\text{out}}=20$, $K=5$, $P=2$, $S=1$)	20 x 28 x 28	20 x 1 x 5 x 5
ReLU	20 x 28 x 28	
MaxPool($K=2$, $S=2$)	20 x 14 x 14	
Conv ($C_{\text{out}}=50$, $K=5$, $P=2$, $S=1$)	50 x 14 x 14	50 x 20 x 5 x 5
ReLU	50 x 14 x 14	
MaxPool($K=2$, $S=2$)	50 x 7 x 7	
Flatten	2450	
Linear (2450 -> 500)	500	2450 x 500
ReLU	500	
Linear (500 -> 10)	10	500 x 10




As we go through the network:

Spatial size **decreases**
(using pooling or strided conv)

Number of channels **increases**
(total “volume” is preserved!)

Lecun et al, “Gradient-based learning applied to document recognition”, 1998

ResNet



Why is it a better idea than flattening the feature map before the classifier (a fully-connect layer)?

Number of parameters.

Spatial dimension of the input.

global average pooling

Batch Normalization

Idea: “Normalize” the outputs of each layer so they have zero mean and unit variance

Why? Helps reduce “internal covariate shift”, improves optimization (hypothesis)

We can normalize a batch of activations like this:

$$\hat{x} = \frac{x - E[x]}{\sqrt{Var[x]}}$$

This is a **differentiable function**, so we can use it as an operator in our networks and backprop through it!

Batch Normalization

Problem: Estimates
depend on minibatch;
can't do this at test-time!

Input: $x \in \mathbb{R}^{N \times D}$

**Learnable scale and
shift parameters:**

$$\gamma, \beta \in \mathbb{R}^D$$

Learning $\gamma = \sigma$, $\beta = \mu$
will recover the identity
function (in expectation)

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

Per-channel
mean, shape is D

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

Per-channel
std, shape is D

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,
Shape is N x D

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

Output,
Shape is N x D

Batch Normalization: Test-Time

Input: $x \in \mathbb{R}^{N \times D}$

$\mu_j =$ (Running) average of values seen during training

Per-channel mean, shape is D

Learnable scale and shift parameters:

$\sigma_j^2 =$ (Running) average of values seen during training

Per-channel std, shape is D

$\gamma, \beta \in \mathbb{R}^D$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Normalized x,
Shape is N x D

During testing batchnorm becomes a linear operator!

Can be fused with the previous fully-connected or conv layer

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

Output,
Shape is N x D

Today's Class

- More about batch normalization layer
- Modern deep convolutional neural networks
- Training deep convolutional neural networks

Batch Normalization for ConvNets


Batch Normalization for
fully-connected networks

$$\begin{array}{l}
 x : N \times D \\
 \text{Normalize} \quad \downarrow \\
 \mu, \sigma : 1 \times D \\
 \gamma, \beta : 1 \times D \\
 y = \frac{(x - \mu)}{\sigma} \gamma + \beta
 \end{array}$$

Batch Normalization for
convolutional networks
(Spatial Batchnorm, BatchNorm2D)

$$\begin{array}{l}
 x : N \times C \times H \times W \\
 \text{Normalize} \quad \downarrow \quad \downarrow \quad \downarrow \\
 \mu, \sigma : 1 \times C \times 1 \times 1 \\
 \gamma, \beta : 1 \times C \times 1 \times 1 \\
 y = \frac{(x - \mu)}{\sigma} \gamma + \beta
 \end{array}$$

Batch Normalization for ConvNets



Batch Normalization for
convolutional networks
(Spatial Batchnorm, BatchNorm2D)

$$x : N \times C \times H \times W$$

Normalizing

$$\mu, \sigma$$


$$: 1 \times C \times 1 \times 1$$

$$\gamma, \beta$$

$$: 1 \times C \times 1 \times 1$$

$$y = \frac{(x - \mu)}{\sigma} \gamma + \beta$$


Batch Normalization




Usually inserted after Fully Connected or Convolutional layers, and before nonlinearity.

$$\hat{x} = \frac{x - E[x]}{\sqrt{Var[x]}}$$


Batch Normalization



- Makes deep networks **much** easier to train!
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Free at test-time: can be fused with conv!




Batch Normalization



Ioffe and Szegedy, "Batch normalization: Accelerating deep network training by reducing internal covariate shift", ICML 2015

- Makes deep networks **much** easier to train!
- Allows higher learning rates, faster convergence
- Networks become more robust to initialization
- Acts as regularization during training
- Free at test-time: can be fused with conv!
- **Not well-understood theoretically (yet)**
- **Behaves differently during training and testing: this is a very common source of bugs!**

Different Normalization Layers




[Wu and He. Group Normalization. ECCV 2018. Best paper honorable mention.]


Devils in the details (mainly for PyTorch)

1. To reliably estimate BN statistics (running mean and average), you need at least 8 samples on each GPU
2. If you don't have enough samples on each GPU
 1. Synchronized BN layers
 2. Group normalization layers
3. Instance normalization layers are useful for some applications: such as style transfer, dense correspondence.


Convolutional Networks



ImageNet Classification Challenge




AlexNet



227 x 227 inputs
5 Convolutional layers
Max pooling
3 fully-connected layers
ReLU nonlinearities

Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

AlexNet



227 x 227 inputs
5 Convolutional layers
Max pooling
3 fully-connected layers
ReLU nonlinearities

Used “Local response normalization”;
Not used anymore

Trained on two GTX 580 GPUs – only
3GB of memory each! Model split
over two GPUs

AlexNet






Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.


Figure copyright Alex Krizhevsky, Ilya Sutskever, and Geoffrey Hinton, 2012. Reproduced with permission.

AlexNet



	Input size		Layer				Output size	
Layer	C	H / W	filters	kernel	stride	pad	C	H / W
conv1	3	227	64	11	4	2	?	


AlexNet



	Input size		Layer				Output size	
Layer	C	H / W	filters	kernel	stride	pad	C	H / W
conv1	3	227	64	11	4	2	64	?

Recall: Output channels = number of filters


AlexNet



	Input size		Layer				Output size	
Layer	C	H / W	filters	kernel	stride	pad	C	H / W
conv1	3	227	64	11	4	2	64	56


$$\begin{aligned}
 \text{Recall: } W' &= \text{ceil}((W - K + 1 + 2P) / S) \\
 &= \text{ceil}((227 - 11 + 1 + 2*2) / 4) \\
 &= \text{ceil}(221/4) = 56
 \end{aligned}$$

AlexNet



	Input size		Layer				Output size		
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)
conv1	3	227	64	11	4	2	64	56	?

AlexNet




	Input size		Layer				Output size		
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)
conv1	3	227	64	11	4	2	64	56	784

$$\begin{aligned}\text{Number of output elements} &= C * H' * W' \\ &= 64 * 56 * 56 = 200,704\end{aligned}$$

$$\text{Bytes per element} = 4 \text{ (for 32-bit floating point)}$$


$$\begin{aligned}\text{KB} &= (\text{number of elements}) * (\text{bytes per elem}) / 1024 \\ &= 200704 * 4 / 1024 \\ &= \mathbf{784}\end{aligned}$$

AlexNet



	Input size		Layer				Output size			
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)
conv1	3	227	64	11	4	2	64	56	784	?

AlexNet




	Input size		Layer				Output size			
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)
conv1	3	227	64	11	4	2	64	56	784	23

$$\begin{aligned}\text{Weight shape} &= C_{\text{out}} \times C_{\text{in}} \times K \times K \\ &= 64 \times 3 \times 11 \times 11\end{aligned}$$

$$\text{Bias shape} = C_{\text{out}} = 64$$


$$\begin{aligned}\text{Number of weights} &= 64 \times 3 \times 11 \times 11 + 64 \\ &= \mathbf{23,296}\end{aligned}$$

AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	?

AlexNet




	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73

Number of floating point operations (multiply+add)


$$\begin{aligned} &= (\text{number of output elements}) * (\text{ops per output elem}) \\ &= (C_{\text{out}} \times H' \times W') * (C_{\text{in}} \times K \times K) \\ &= (64 * 56 * 56) * (3 * 11 * 11) \\ &= 200,704 * 363 \\ &= \mathbf{72,855,552} \end{aligned}$$

AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	?				

AlexNet




Layer	Input size		Layer				Output size				
	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27			

For pooling layer:

#output channels = #input channels = 64

$$\begin{aligned}
 W' &= \text{ceil}((W - K + 1) / S) \\
 &= \text{ceil}(54 / 2) = \text{ceil}(27) = \mathbf{27}
 \end{aligned}$$

AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	?	

$$\# \text{output elems} = C_{\text{out}} \times H' \times W'$$


$$\text{Bytes per elem} = 4$$

$$\text{KB} = C_{\text{out}} * H' * W' * 4 / 1024$$

$$= 64 * 27 * 27 * 4 / 1024$$

$$= \mathbf{182.25}$$


AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	?

Pooling layers have no learnable parameters!

AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0

Floating-point ops for pooling layer

= (number of output positions) * (flops per output position)


= $(C_{\text{out}} * H' * W') * (K * K)$

= $(64 * 27 * 27) * (3 * 3)$

= 419,904

= **0.4 MFLOP**


AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0
conv2	64	27	192	5	1	2	192	27	547	307	224
pool2	192	27		3	2	0	192	13	127	0	0
conv3	192	13	384	3	1	1	384	13	254	664	112
conv4	384	13	256	3	1	1	256	13	169	885	145
conv5	256	13	256	3	1	1	256	13	169	590	100
pool5	256	13		3	2	0	256	6	36	0	0
flatten	256	6					9216		36	0	0

$$\begin{aligned}
 \text{Flatten output size} &= C_{in} \times H \times W \\
 &= 256 * 6 * 6 \\
 &= \mathbf{9216}
 \end{aligned}$$

AlexNet




	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0
conv2	64	27	192	5	1	2	192	27	547	307	224
pool2	192	27		3	2	0	192	13	127	0	0
conv3	192	13	384	3	1	1	384	13	254	664	112
conv4	384	13	256	3	1	1	256	13	169	885	145
conv5	256	13	256	3	1	1	256	13	169	590	100
pool5	256	13		3	2	0	256	6	36	0	0
flatten	256	6					9216		36	0	0
fc6	9216		4096				4096		16	37,749	38

$$\begin{aligned}
 \text{FC params} &= C_{\text{in}} * C_{\text{out}} + C_{\text{out}} \\
 &= 9216 * 4096 + 4096 \\
 &= 37,725,832
 \end{aligned}$$

$$\begin{aligned}
 \text{FC flops} &= C_{\text{in}} * C_{\text{out}} \\
 &= 9216 * 4096 \\
 &= 37,748,736
 \end{aligned}$$


AlexNet



	Input size		Layer				Output size				
Layer	C	H / W	filters	kernel	stride	pad	C	H / W	memory (KB)	params (k)	flop (M)
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0
conv2	64	27	192	5	1	2	192	27	547	307	224
pool2	192	27		3	2	0	192	13	127	0	0
conv3	192	13	384	3	1	1	384	13	254	664	112
conv4	384	13	256	3	1	1	256	13	169	885	145
conv5	256	13	256	3	1	1	256	13	169	590	100
pool5	256	13		3	2	0	256	6	36	0	0
flatten	256	6					9216		36	0	0
fc6	9216		4096				4096		16	37,749	38
fc7	4096		4096				4096		16	16,777	17
fc8	4096		1000				1000		4	4,096	4


AlexNet

How to choose this?
Trial and error =(



Layer	Input size		Layer				Output size		memory (KB)	params (k)	flop (M)
	C	H / W	filters	kernel	stride	pad	C	H / W			
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0
conv2	64	27	192	5	1	2	192	27	547	307	224
pool2	192	27		3	2	0	192	13	127	0	0
conv3	192	13	384	3	1	1	384	13	254	664	112
conv4	384	13	256	3	1	1	256	13	169	885	145
conv5	256	13	256	3	1	1	256	13	169	590	100
pool5	256	13		3	2	0	256	6	36	0	0
flatten	256	6					9216		36	0	0
fc6	9216		4096				4096		16	37,749	38
fc7	4096		4096				4096		16	16,777	17
fc8	4096		1000				1000		4	4,096	4


AlexNet



Interesting trends here!


Layer	Input size		Layer				Output size		memory (KB)	params (k)	flop (M)
	C	H / W	filters	kernel	stride	pad	C	H / W			
conv1	3	227	64	11	4	2	64	56	784	23	73
pool1	64	56		3	2	0	64	27	182	0	0
conv2	64	27	192	5	1	2	192	27	547	307	224
pool2	192	27		3	2	0	192	13	127	0	0
conv3	192	13	384	3	1	1	384	13	254	664	112
conv4	384	13	256	3	1	1	256	13	169	885	145
conv5	256	13	256	3	1	1	256	13	169	590	100
pool5	256	13		3	2	0	256	6	36	0	0
flatten	256	6					9216		36	0	0
fc6	9216		4096				4096		16	37,749	38
fc7	4096		4096				4096		16	16,777	17
fc8	4096		1000				1000		4	4,096	4

AlexNet




Most of the **memory usage** is in the early convolution layers

Memory (KB)




Nearly all **parameters** are in the fully-connected layers

Params (K)




Most **floating-point ops** occur in the convolution layers


MFLOP



ImageNet Classification Challenge



ImageNet Classification Challenge



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels



Simonyan and Zissermann, "Very Deep Convolutional Networks for Large-Scale Image Recognition", ICLR 2015

VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Network has 5 convolutional **stages**:

Stage 1: conv-conv-pool

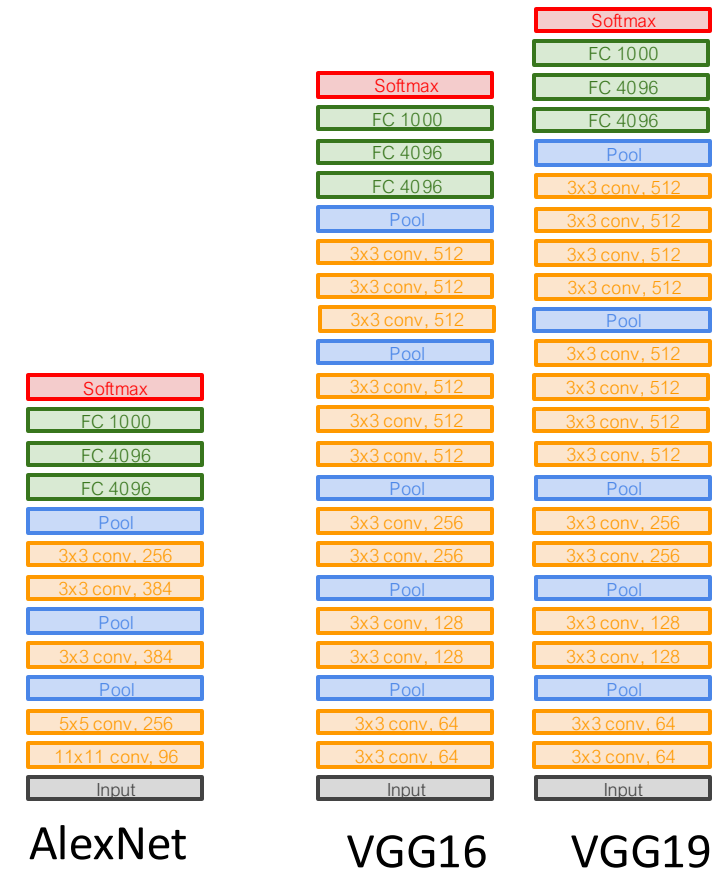
Stage 2: conv-conv-pool

Stage 3: conv-conv-pool

Stage 4: conv-conv-conv-[conv]-pool

Stage 5: conv-conv-conv-[conv]-pool

(VGG-19 has 4 conv in stages 4 and 5)



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Option 1:

Conv(5x5, C → C)

Params: $25C^2$

FLOPs: $25C^2HW$



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Option 1:

Conv(5x5, C → C)

Params: $25C^2$

FLOPs: $25C^2HW$


Option 2:

Conv(3x3, C → C)

Conv(3x3, C → C)

Params: $18C^2$

FLOPs: $18C^2HW$



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Option 1:

Conv(5x5, C → C)

Params: $25C^2$

FLOPs: $25C^2HW$

Option 2:

Conv(3x3, C → C)

Conv(3x3, C → C)

Params: $18C^2$

FLOPs: $18C^2HW$

Two 3x3 conv has same receptive field as a single 5x5 conv, but has fewer parameters and takes less computation!



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels


Input: $C \times 2H \times 2W$

Layer: Conv(3x3, C->C)

Memory: 4HWC

Params: $9C^2$

FLOPs: $36HWC^2$



Simonyan and Zissermann, "Very Deep Convolutional Networks for Large-Scale Image Recognition", ICLR 2015

VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Input: $C \times 2H \times 2W$

Layer: Conv(3x3, $C \rightarrow C$)

Memory: $4HWC$

Params: $9C^2$

FLOPs: $36HWC^2$

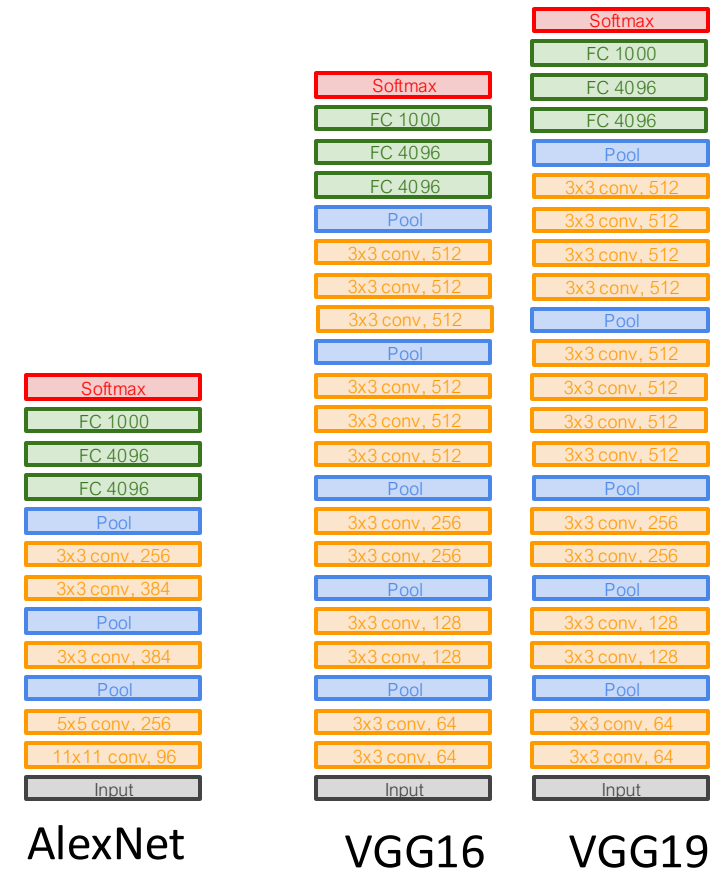
Input: $2C \times H \times W$

Conv(3x3, $2C \rightarrow 2C$)

Memory: $2HWC$

Params: $36C^2$

FLOPs: $36HWC^2$



VGG: Deeper Networks, Regular Design

VGG Design rules:

All conv are 3x3 stride 1 pad 1

All max pool are 2x2 stride 2

After pool, double #channels

Input: $C \times 2H \times 2W$

Layer: Conv(3x3, $C \rightarrow C$)

Memory: $4HWC$

Params: $9C^2$

FLOPs: $36HWC^2$

Input: $2C \times H \times W$

Conv(3x3, $2C \rightarrow 2C$)

Memory: $2HWC$


Params: $36C^2$

FLOPs: $36HWC^2$


Conv layers at each spatial resolution take the same amount of computation!



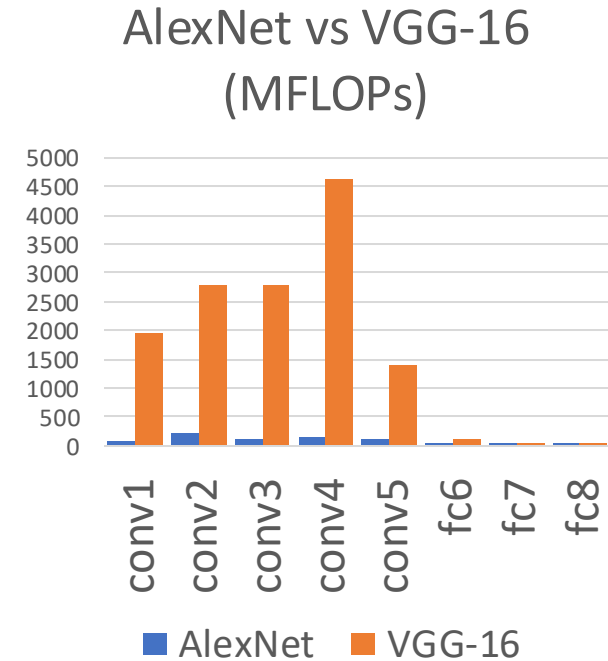
AlexNet vs VGG-16: Much Bigger!



AlexNet total: 1.9 MB
VGG-16 total: 48.6 MB (25x)




AlexNet total: 61M
VGG-16 total: 138M (2.3x)




AlexNet total: 0.7 GFLOP
VGG-16 total: 13.6 GFLOP (19.4x)

ImageNet Classification Challenge



ImageNet Classification Challenge



Residual Networks

Once we have Batch Normalization, we can train networks with 10+ layers. What happens as we go deeper?


He et al, "Deep Residual Learning for Image Recognition", CVPR 2016

Residual Networks

Once we have Batch Normalization, we can train networks with 10+ layers. What happens as we go deeper?

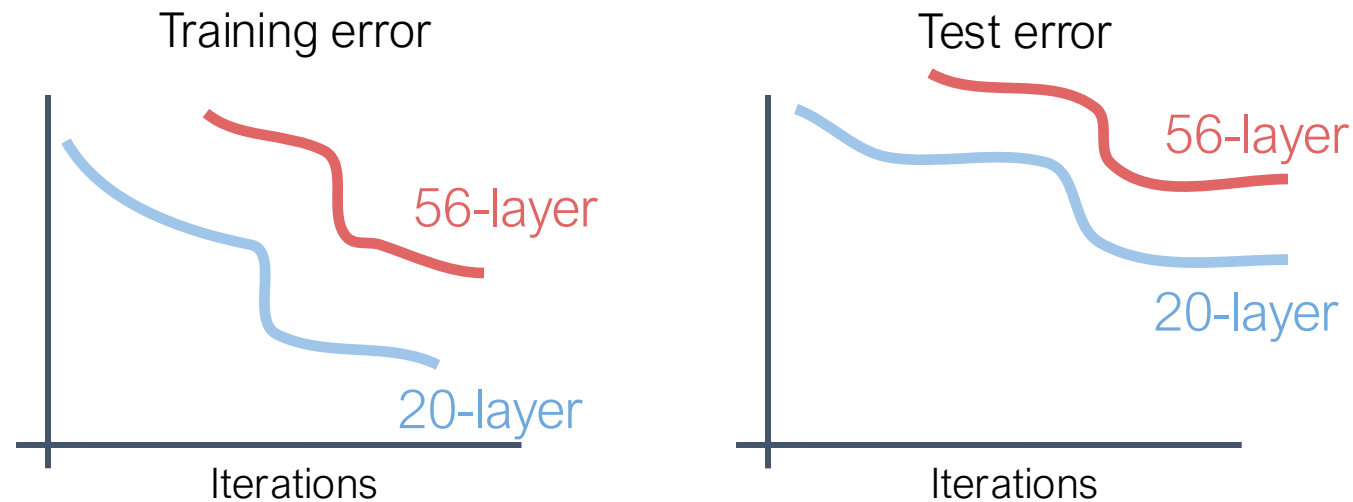
Deeper model does worse than shallow model!

Initial guess: Deep model is **overfitting** since it is much bigger than the other model



Residual Networks

Once we have Batch Normalization, we can train networks with 10+ layers. What happens as we go deeper?



In fact the deep model seems to be **underfitting** since it also performs worse than the shallow model on the training set! It is actually **underfitting**

He et al, "Deep Residual Learning for Image Recognition", CVPR 2016

Residual Networks

A deeper model can emulate a shallower model: copy layers from shallower model, set extra layers to identity

Thus deeper models should do at least as good as shallow models

Hypothesis: This is an optimization problem. Deeper models are harder to optimize, and in particular don't learn identity functions to emulate shallow models

Residual Networks

A deeper model can emulate a shallower model: copy layers from shallower model, set extra layers to identity



Thus deeper models should do at least as good as shallow models

Hypothesis: This is an optimization problem. Deeper models are harder to optimize, and in particular don't learn identity functions to emulate shallow models

Solution: Change the network so learning identity functions with extra layers is easy!

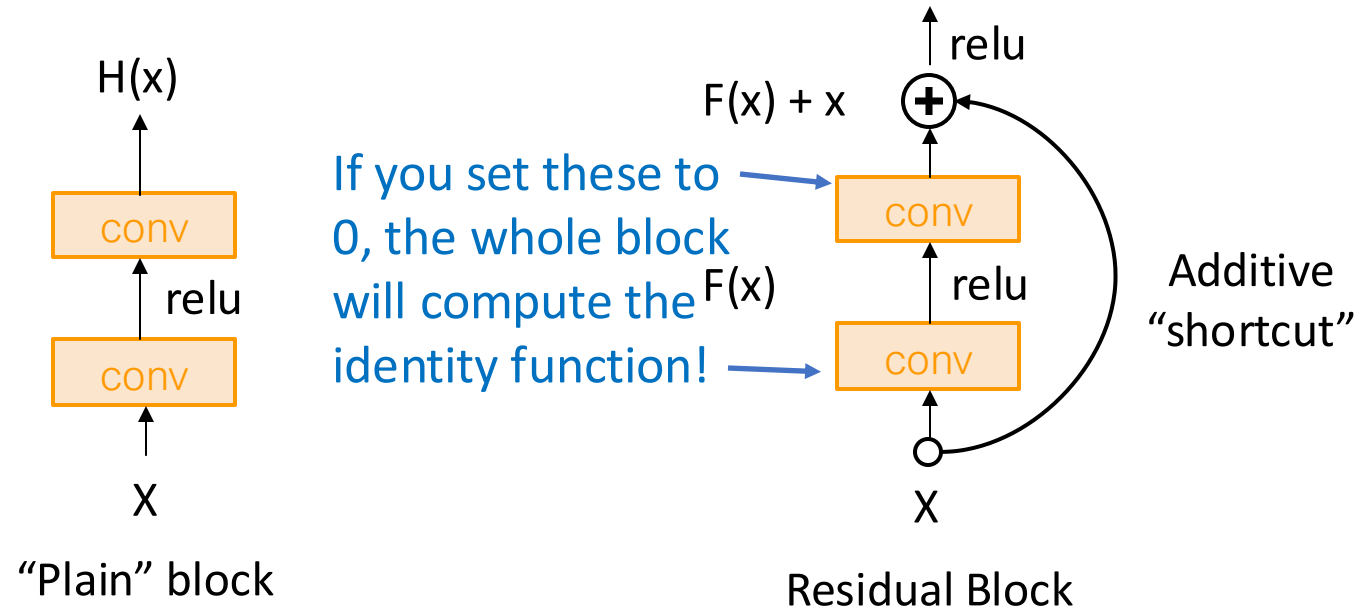
Residual Networks

Solution: Change the network so learning identity functions with extra layers is easy!



Residual Networks

Solution: Change the network so learning identity functions with extra layers is easy!

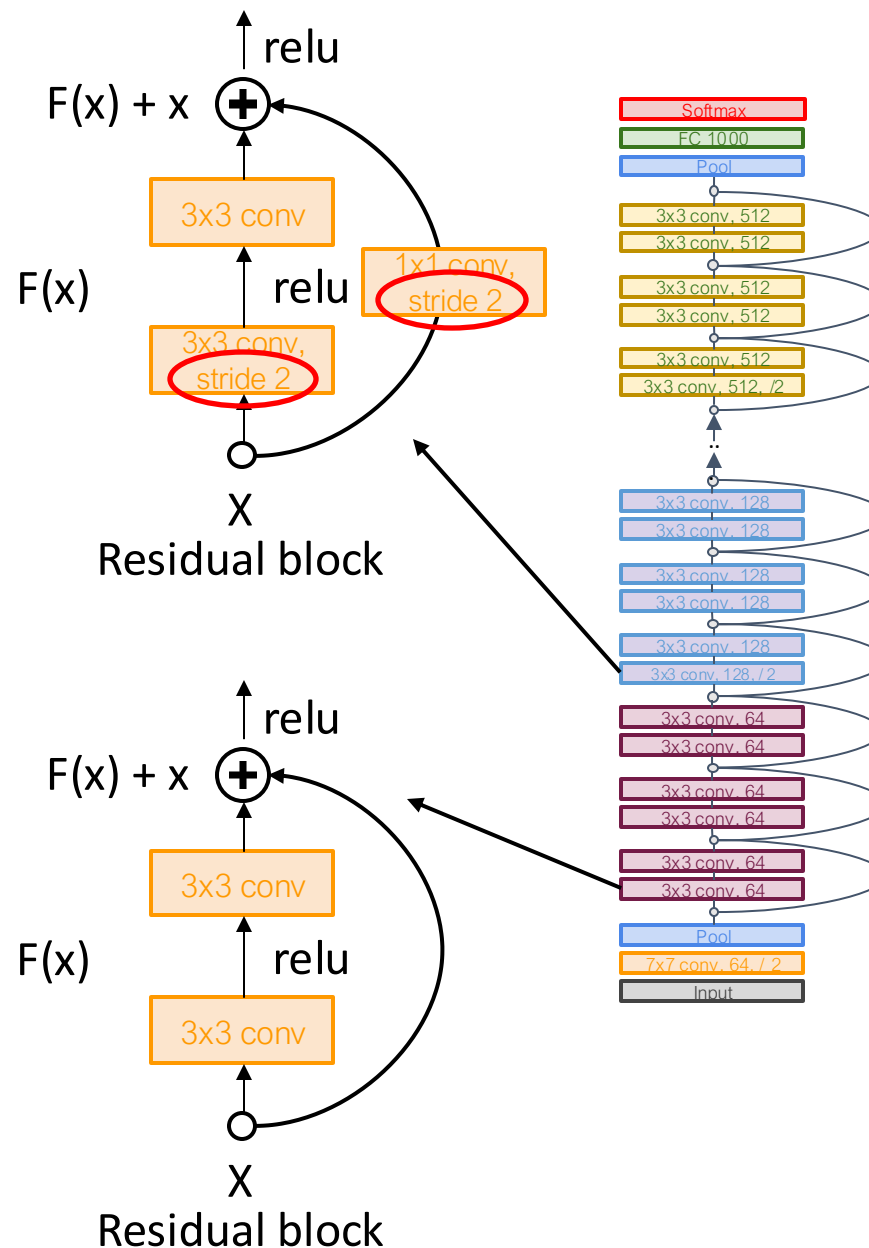


Residual Networks

A residual network is a stack of many residual blocks


Regular design, like VGG:
each residual block has two
3x3 conv

Network is divided into **stages**: the first block of each stage halves the resolution (with stride-2 conv) and doubles the number of channels




He et al, "Deep Residual Learning for Image Recognition", CVPR 2016


ImageNet Classification Challenge



ImageNet Classification Challenge



ImageNet Classification Challenge



Tiny Networks for Mobile Devices

Object Detection



Photo by Juanedc (CC BY 2.0)

Face Attributes



Google Doodle by Sarah Harrison

Finegrain Classification




Photo by HarshLight (CC BY 2.0)

Landmark Recognition




Photo by Sharon VanderKaay (CC BY 2.0)



[Howard et al., MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications. arXiv 2017]

Group-based Convolution



Input: $C_{in} \times H \times W$

Hyperparameters:


- **Kernel size:** $K_H \times K_W$
- **Number filters:** C_{out}
- **Padding:** P
- **Stride:** S
- **Groups:** G

Weight matrix: $C_{out}/G \times C_{in}/G \times K_H \times K_W \times G$

Bias vector: C_{out}/G


FLOPS: $C_{out}/G \times C_{in}/G \times K_H \times K_W \times G \times H \times W$

MobileNet



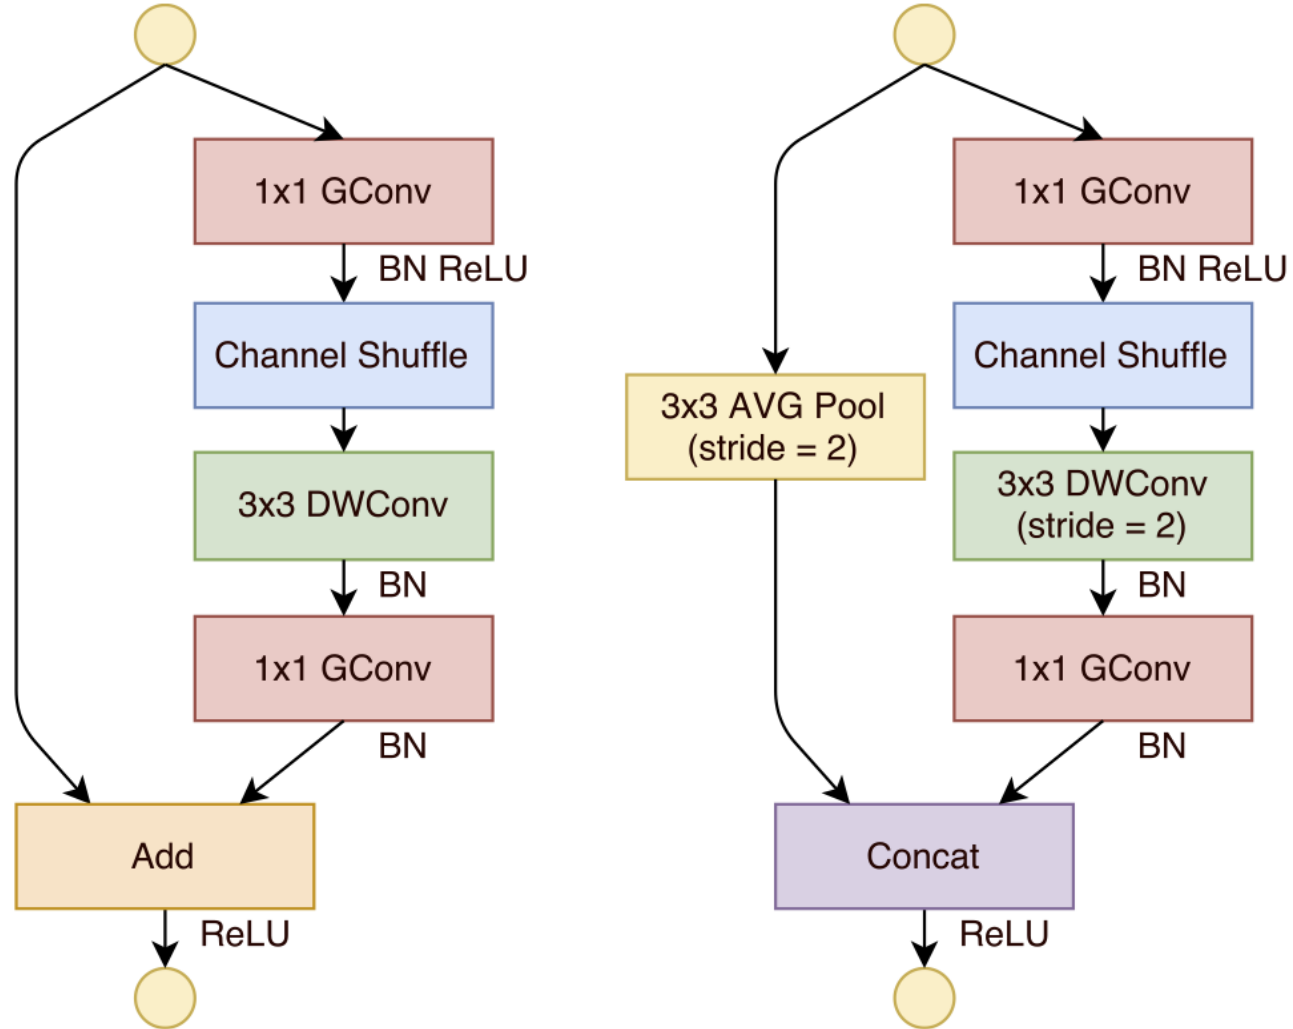
Computation reduction: $9C^2HW / (9CHW + C^2HW) = 9C / (9 + C)$

ShuffleNet



[Zhang et al., ShuffleNet: An Extremely Efficient Convolutional Neural Network for Mobile Devices. CVPR 2018]

ShuffleNet Units



Training Convolutional Networks

1. Download big datasets
2. Design CNN architecture
3. Initialize Weights
4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient
 3. Update Weights
5. Apply trained model to task

Training Convolutional Networks

1. Download big datasets
2. Design CNN architecture
- 3. Initialize Weights**
4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient
 3. Update Weights
5. Apply trained model to task

Training Convolutional Networks

1. Download big datasets
 2. Design CNN architecture
 3. Initialize Weights
 4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient
 3. Update Weights
 5. Apply trained model to task
- If the model is big, won't we overfit?

Regularizing CNNs: Weight Decay

$$L_{reg} = \frac{1}{2} \sum_{\ell} \|W_{\ell}\|^2 \quad \frac{\partial L_{reg}}{\partial W_{\ell}} = W_{\ell}$$

Add L2 regularization term L_{reg} to the loss penalizing large weight matrices

Usually don't regularize bias terms, or BatchNorm scale / shift params


*Technical note: Adding an explicit term to the loss is "L2 Regularization"; "Weight decay" adds a term to the gradient. They are equivalent for SGD, but not quite the same for other optimizers like Adam

Regularizing CNNs: Data Augmentation

Hippo




Hippo?



Horizontal
Flip

Hippo?



Color
Jitter

Hippo?





Image
Cropping

Regularizing CNNs: Data Augmentation

Apply random transformations to input images during training
Artificially “inflate” the size of your dataset



Training Convolutional Networks


1. Download big datasets
 2. Design CNN architecture
 3. Initialize Weights
 4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient
 3. Update Weights
 5. Apply trained model to task
- If the model is big, won't we overfit?

Training Convolutional Networks

1. Download big datasets
 2. Design CNN architecture
 3. Initialize Weights
 4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient
 3. Update Weights
 5. Apply trained model to task
- What if
we can't
find one?


Transfer Learning: Feature Extraction

1. Train on ImageNet




Transfer Learning: Feature Extraction

1. Train on ImageNet



2. CNN as feature extractor




Use your small dataset to train a **linear classifier** on top of pretrained CNN features

Transfer Learning: Fine-Tuning

1. Train on ImageNet



2. CNN as feature extractor



Remove last layer

Freeze these

3. Bigger dataset: **Fine-Tuning**




Reinitialize last layer and continue training whole network on your dataset

Transfer Learning: Fine-Tuning

1. Train on ImageNet



2. CNN as feature extractor



Remove last layer

Freeze these

3. Bigger dataset: Fine-Tuning




Reinitialize last layer and continue training whole network on your dataset

Some tricks:

- Train with feature extraction first before fine-tuning
- Lower the learning rate: use $\sim 1/10$ of LR used in original training
- Sometimes freeze lower layers to save computation


Recap: Convolutional Networks

Fully-Connected Layers




$$y = Wx + b$$

Activation Function




$$y = \max(0, x)$$

Convolution Layers




Pooling Layers



Normalization

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \varepsilon}}$$

Recap: CNN Architectures



Recap: Training CNNs

1. Download big datasets Transfer Learning
2. Design CNN architecture
3. Initialize Weights Xavier / MSRA Init
4. For $t = 1$ to T :
 1. Form minibatch
 2. Compute loss + gradient Regularization + Data
 3. Update Weights Augmentation
5. Apply trained model to task

Next Class

More about
Convolutional Neural Networks